

Problem Set 4

published: 12 May 2021

discussion: 17 May 2021

Data Compression With Deep Probabilistic Models

Prof. Robert Bamler, University of Tuebingen

Course material available at <https://robamler.github.io/teaching/compress21/>

Notes:

- This problem set is deliberately kept short with no coding problem because we'll use most of the time in the tutorial on May 17 for a final matchmaking round for the group projects.
- Problem 4.2 introduces the concepts of joint and conditional entropies. You will derive important and non-obvious relations between these concepts that will be instrumental throughout the rest of the course. These relations will guide both the choice of machine-learning models as well as the design of coding algorithms, and they will be critical for the theories of lossy compression and noisy communication.

Problem 4.1: Joint and Conditional Information Content

On the last problem set, we analyzed the bit rate of a lossless compression algorithm that is optimal w.r.t. some probabilistic model p_{model} of the data source. We saw that, up to usually negligible rounding effects, the bit rate $R(\mathbf{x})$ for any message \mathbf{x} under such a code is given by the *information content* of \mathbf{x} under the model, i.e., $R(\mathbf{x}) \simeq -\log p_{\text{model}}(\mathbf{x})$. Thus, we can understand the information content as the number of bits that are essential to a message, regardless of the representation in which the message is actually stored or transmitted.

Now that we have a more formal notion of probability theory, we can analyze how each symbol in the message contributes to the information content. In the general case, it is best to regard the symbols in a message as *random variables* X_1, X_2, \dots, X_k . But for this problem, we'll just look at two random variables X and Y . The generalization to more than two random variables is straight forward. In this problem, we assume that X and Y are both *discrete* random variables, since the information content is not defined for continuous random variables. We denote the probabilistic model simply as P (dropping the subscript "model").

- (a) **Joint Information Content:** The joint information content for X and Y is the information content for the tuple (X, Y) , which can be considered as a random variable in itself ($\omega \mapsto (X(\omega), Y(\omega))$). Thus, for some specific values x and y , the information content is

$$\begin{aligned} -\log P((X, Y) = (x, y)) &= -\log P(X = x, Y = y) \\ &= -\log P(\{\omega \in \Omega : X(\omega) = x \wedge Y(\omega) = y\}). \end{aligned} \tag{1}$$

Show that the information content of $(X, Y) = (x, y)$ is not smaller than the information content of $X = x$ alone and not smaller than the information content of $Y = y$ alone (the information content of $X = x$ is $-\log P(X = x) = -\log P(\{\omega \in \Omega : X(\omega) = x\})$). Thus, adding an additional symbol can't decrease the information content.

- (b) **Marginal and Conditional Information Content:** The information content of $X = x$ alone, $-\log P(X = x)$, is also called the *marginal* information content. We further define the *conditional* information content of $Y = y$ given $X = x$ as $-\log P(Y = y|X = x)$. Using the definition of the conditional probability from the lecture, $P(Y = y|X = x) := P(X = x, Y = y)/P(X = x)$, derive the chain rule of information content, which states that the joint information content of $(X, Y) = (x, y)$ is the sum of the marginal information content of $X = x$ and the conditional information content of $Y = y$ given $X = x$. What does this mean for the bit rate of an autoregressive model?
- (c) **Nonadditivity of Marginal Information Content:** In general, the joint information content of $(X, Y) = (x, y)$ is *not* the sum of the two marginal information contents of $X = x$ and $Y = y$. In fact, show (by providing a simple example for both cases) that the sum of the two marginal information contents can be both lower and higher than the joint information content.
- (d) **Information Content for Statistically Independent Random Variables:** Now assume that X and Y are statistically independent, i.e., $P(X, Y) = P(X)P(Y)$ (which is shorthand for $P(X = x, Y = y) = P(X = x)P(Y = y) \forall x, y$). How does the statement from part (c) change in this case.

Problem 4.2: Joint and Conditional Entropy

As we defined in the lecture, the entropy $H_P(X)$ of a random variable X is its *expected information content*, i.e., $H_P(X) = \mathbb{E}_P[-\log P(X)]$. Similar to Problem 4.1, let's now understand how entropies of two random variables X and Y interact. We will again assume that X and Y are *discrete* random variables since the entropy is not defined for continuous random variables (only a *differential* entropy is defined for these).

- (a) **Joint Entropy:** The joint entropy of X and Y , is the entropy of the tuple (X, Y) . We will explicitly denote it as $H_P((X, Y))$ (with double braces) to highlight the distinction from the cross entropy.¹ Show, by applying what you've shown in Problem 4.1 (a), that $H_P((X, Y)) \geq H_P(X)$ and that $H_P((X, Y)) \geq H_P(Y)$.
- (b) **Marginal and Conditional Entropy:** The entropy of X alone, $H_P(X)$, is also called the *marginal* entropy. We further define two kinds of conditional entropies:

¹This is not really standard notation, the literature is inconsistent in the notation here; you may find $H(X, Y)$ to refer to either the cross entropy or the joint entropy in the literature.

- $H_P(Y|X = x)$ denotes the conditional entropy of Y if we know that X takes a specific value x . In other words, $H_P(Y|X = x)$ is the entropy of the distribution $P(Y|X = x)$, interpreted as a distribution over values of Y . It is thus given by

$$\begin{aligned} H_P(Y|X = x) &= \mathbb{E}_P[-\log P(Y|X = x)] \\ &= -\sum_y P(Y = y | X = x) \log P(Y = y | X = x) \end{aligned} \quad (2)$$

Show (by providing an example for both cases) that $H_P(Y|X = x)$ can be both larger and smaller than $H_P(Y)$.

- $H_P(Y|X)$ denotes the expectation value of $H_P(Y|X = x)$, where the expectation is taken over x . Thus,

$$\begin{aligned} H_P(Y|X) &= \sum_x P(X = x) H_P(Y|X = x) \\ &= -\sum_x P(X = x) \left[\sum_y P(Y = y | X = x) \log P(Y = y | X = x) \right] \\ &= -\sum_{x,y} P(X = x, Y = y) \log P(Y = y | X = x) \end{aligned} \quad (3)$$

Derive the chain rule of the entropy:

$$H_P((X, Y)) = H_P(X) + H_P(Y|X) = H_P(Y) + H_P(X|Y). \quad (4)$$

We will show on the next problem set that $H_P(Y|X) \leq H_P(Y)$ (the gap between the two is called the “mutual information between X and Y ”). Therefore, while conditioning on a *specific* $X = x$ may increase the conditional entropy $H_P(Y|X = x)$ (see above), *in expectation*, conditioning can only decrease the entropy (or keep it unchanged at worst).

- (c) **Entropy of Statistically Independent Random Variables:** What is the joint entropy $H_P((X, Y))$ and the marginal entropy $H_P(Y|X = x)$ and $H_P(Y|X)$ if the two random variables X and Y are statistically independent, i.e., if $P(X, Y) = P(X)P(Y)$?

Don't forget to provide anonymous feedback to this problem set in the corresponding poll on [moodle](#).