Lecture 2, Part 1:

Theoretical Bounds for Lossless Compression

Robert Bamler · Summer Term of 2023

These slides are part of the course *"Data Compression With and Without Deep Probabilistic Models"* taught at University of Tübingen. More course materials—including video recordings, lecture notes, and problem sets with solutions—are publicly available at `https://robamler.github.io/teaching/compress23/`.

Recap: Symbol Codes

- *alphabet* $\mathfrak{X}$ (discrete set) with *probabilities* $p(x)$ for all *symbols* $x \in \mathfrak{X}$
- *message* $\mathbf{x} = (x_1, x_2, \ldots, x_{k(\mathbf{x})}) \in \mathfrak{X}^*$
- *code book* C maps any $x \in \mathfrak{X}$ to its *code word* $C(x) \in \{0, \ldots, B-1\}^*$ (usually: $B = 2$)
  - induces a *symbol code* $C^*\colon \mathfrak{X}^* \to \{0, \ldots, B-1\}^*$ by concatenation (without delimiters): $C^*(\mathbf{x}) := C(x_1) \,\|\, C(x_2) \,\|\, \ldots \,\|\, C(x_{k(\mathbf{x})})$
- properties of symbol codes:
  - *unique decodability:* C^* is injective
  - *prefix code:* no code word $C(x)$ is a prefix of another code word $C(x')$ with $x' \neq x$
  - C is a prefix code $\Rightarrow$ C is uniquely decodable (but reverse is in general not true)
- *expected code word length* $L_C := \sum_{x \in \mathfrak{X}} p(x)\,|C(x)|$
- *Huffman coding* generates an optimal symbol code (that minimizes L_C) for a given p

Theoretical Bounds for Lossless Compression

- **Goal of this lecture:** Source Coding Theorem [Shannon, 1948]
  - Relates L_C to the so-called *entropy* $H_B[p]$ (which we'll define later today).
  - **The Bad News:** no uniquely decodable B-ary symbol code C can have $L_C < H_B[p]$.
  - **The Good News:** $\forall p$, one can make L_C close to $H_B[p]$ with less than 1 bit per symbol overhead.
- **Step 1:** proof bound on code word lengths, independently from p (KM-Theorem)
- **Step 2:** proof bound on *expected* code word length for a given model p
- **Credits:** Our proof follows: `https://youtu.be/TODO`

The Kraft-McMillan Theorem [Kraft, 1949; McMillan, 1956]

EBERHARD KARLS UNIVERSITÄT TÜBINGEN

(a) $\forall$ B-ary uniquely decodable symbol codes over some discrete alphabet $\mathfrak{X}$:

$$\sum_{x\in\mathfrak{X}} \frac{1}{B^{|C(x)|}} \leq 1 \qquad (\text{“Kraft inequality”}). \tag{1}$$

Interpretation: we have a finite budget of “shortness” for code words:

- interpret $\frac{1}{B^{|C(x)|}}$ as the “shortness” of code word $C(x)$;
- the sum of all “shortnesses” must not exceed 1;
- if we shorten one code word then we may have to make another code word longer so that we don't exceed our “shortness budget”.

(b) $\forall$ functions $\ell : \mathfrak{X} \to \mathbb{N}$ that satisfy the Kraft inequality (i.e., $\sum_{x\in\mathfrak{X}} \frac{1}{B^{\ell(x)}} \leq 1$):

$$\exists\ B\text{-ary prefix code } C_\ell \text{ with } |C_\ell(x)| = \ell(x)\ \ \forall x\in\mathfrak{X}.$$

Corollary: $\forall$ uniquely decodable B-ary symbol codes C:
$\exists$ a B-ary *prefix* code C' with same code word lengths (i.e., $|C'(x)| = |C(x)|\ \forall x\in\mathfrak{X}$)

Lemma

EBERHARD KARLS UNIVERSITÄT TÜBINGEN

- let: $\begin{cases} C \text{ be a } B\text{-ary uniquely decodable symbol code over } \mathfrak{X}; \\ s \in \mathbb{N}_0; \\ Y_s := \{\mathbf{x} \in \mathfrak{X}^* \text{ with } |C^*(\mathbf{x})| = s\}. \end{cases}$
- then: $|Y_s| \leq B^s$.

Proof:

Proof of Part (a) of KM Theorem

EBERHARD KARLS UNIVERSITÄT TÜBINGEN

Claim (reminder): C is uniquely decodable $\implies \sum_{x\in\mathfrak{X}} \frac{1}{B^{|C(x)|}} \leq 1$.

(i) if $\mathfrak{X}$ is finite:

(ii) if $\mathfrak{X}$ is countably infinite:

Proof of Part (b) of KM Theorem

Claim (reminder): $\sum_{x\in\mathfrak{X}} \frac{1}{B^{\ell(x)}} \leq 1 \implies \exists\ B$-ary prefix code C_ℓ with $|C_\ell(x)| = \ell(x)\ \forall x \in \mathfrak{X}$.

Constructive proof: we show existence of C by showing how it can be obtained.

Claim: The resulting code book C_ℓ is prefix free (proof: Problem 2.1).

Example: Simplified Game of Monopoly (SGoM)

x	$\ell(x)$	$C_\ell(x)$
2	3	
3	2	
4	2	
5	2	
6	3	

- Check Kraft inequality for $B = 2$:
- **Question:** how should we choose $\ell : \mathfrak{X} \to \mathbb{N}$ for a given probabilistic model p?
  - **optimally:** via Huffman coding
  - **near-optimally:** via *information content* (next part).

Outlook

- **Problem Set 2:**
  - complete proof of part (b) of KM-Theorem
  - implement Huffman *decoder* in Python
- **Next part:**
  - theoretical bounds on the expected code word length L_C ("The Bad News" & "The Good News")
  - theoretical bounds *beyond symbol codes:* Source Coding Theorem

Lecture 2, Part 2:

The Source Coding Theorem

Robert Bamler · Summer Term of 2023

These slides are part of the course *"Data Compression With and Without Deep Probabilistic Models"* taught at University of Tübingen. More course materials—including video recordings, lecture notes, and problem sets with solutions—are publicly available at `https://robamler.github.io/teaching/compress23/`.

Recap: Kraft-McMillan (KM) Theorem

(a) $\forall$ B-ary uniquely decodable symbol codes over some discrete alphabet $\mathfrak{X}$:

$$\sum_{x\in\mathfrak{X}} \frac{1}{B^{|C(x)|}} \leq 1 \qquad (\text{"Kraft inequality"}). \tag{1}$$

(b) $\forall$ functions $\ell : \mathfrak{X} \to \mathbb{N}$ that satisfy the Kraft inequality (i.e., $\sum_{x\in\mathfrak{X}} \frac{1}{B^{\ell(x)}} \leq 1$):

$$\exists\ B\text{-ary prefix code } C_\ell \text{ with } |C_\ell(x)| = \ell(x)\ \forall x \in \mathfrak{X}.$$

- **Question:** how should we choose $\ell : \mathfrak{X} \to \mathbb{N}$ for a given probabilistic model p?
  - **optimally:** via Huffman coding (problem: no closed-form solution)
  - **near-optimally (this part):** via *information content*
    spoiler: $\ell_S(x) := \lceil -\log_B p(x) \rceil$

Optimal Choice of ℓ

- **Constrained optimization problem:** $(\star)$
  - minimize: $L_{C_\ell} = \sum_{x\in\mathfrak{X}} p(x)\,|C_\ell(x)| = \sum_{x\in\mathfrak{X}} p(x)\,\ell(x)$
  - constraints: (i) $\sum_{x\in\mathfrak{X}} \frac{1}{B^{\ell(x)}} \leq 1$; (ii) $\ell(x) \in \mathbb{N}\ \ \forall x \in \mathfrak{X}$.
- **Idea:** relax constraint (ii): $(\square)$
  - minimize: $L_\ell := \sum_{x\in\mathfrak{X}} p(x)\,\ell(x)$
  - constraints: (i) $\sum_{x\in\mathfrak{X}} \frac{1}{B^{\ell(x)}} \leq 1$; (ii') $\ell(x) \in \mathbb{R}_{>0}\ \ \forall x \in \mathfrak{X}$.

  $\Rightarrow$ yields *lower bound*: solution L_ℓ of $(\square)$ $\leq$ solution L_{C_ℓ} of $(\star)$
- **Observation:** solution of $(\square)$ satisfies: (i') $\sum_{x\in\mathfrak{X}} \frac{1}{B^{\ell(x)}} = 1$.
  - Enforce via Lagrange multiplier λ:
    find stationary point of $\mathcal{L}_{\ell,\lambda} := L_\ell + \lambda\big(\sum_{x\in\mathfrak{X}} \frac{1}{B^{\ell(x)}} - 1\big)$ w.r.t. $\lambda \in \mathbb{R}$ and all $\ell(x) \in \mathbb{R}_{\geq 0}\ \forall x \in \mathfrak{X}$.

Lower Bound on Expected Code Word Length L_C

EBERHARD KARLS UNIVERSITÄT TÜBINGEN

- **Solution of relaxed optimization problem ($\square$):** $\ell(x) = \underbrace{-\log_B p(x)}_{\substack{\text{“information content of the symbol } x\text{”} \\ \text{(under model } p \text{ and to base } B)}}$
  - $L_\ell = \sum_{x\in\mathfrak{X}} p(x)\,\ell(x) = \underbrace{-\sum_{x\in\mathfrak{X}} p(x)\log_B p(x)}_{=:H_B[p] \text{ (“entropy”)}}$
- Let's now restore the constraints from ($\star$), i.e., $\ell : \mathfrak{X} \to \mathbb{N}$ must be *integer valued.*
  - **Recall:** solution L_{C_ℓ} of ($\star$) $\geq$ solution L_ℓ of ($\square$)
  - Thus, for all *integer valued* ℓ that satisfy Kraft inequality: $L_{C_\ell} \geq H_B[p]$
- By part (a) of the KM-Theorem:

$$\boxed{\begin{array}{c}\textbf{lower bound on the expected code word length } L_C \\ \textbf{of any uniquely decodable } B\textbf{-ary symbol code } C\textbf{:} \\ L_C \geq H_B[p]\end{array}}$$

Shannon Coding [Shannon, 1948]

EBERHARD KARLS UNIVERSITÄT TÜBINGEN

- **Last slide:**
  - Lower bound for uniquely decodable B-ary symbol code: $L_C \geq H_B[p] = -\sum_{x\in\mathfrak{X}} p(x)\log_B p(x)$
  - We would achieve equality ($L_C = H_B[p]$) if we were able to set $\ell(x) = \underbrace{-\log_B p(x)}_{\notin\mathbb{N} \text{ (in general)}}$ $\forall x \in \mathfrak{X}$.
- **Question:** How closely can we approach this bound?
- **Idea:** choose $\ell_S : \mathfrak{X} \to \mathbb{N}$ as follows: $\ell_S(x) = \lceil -\log_B p(x) \rceil$
  - Satisfies Kraft inequality: $\sum_{x\in\mathfrak{X}} B^{-\ell_S(x)} = \sum_{x\in\mathfrak{X}} B^{-\lceil -\log_B p(x)\rceil} \leq \sum_{x\in\mathfrak{X}} B^{\log_B p(x)} = \sum_{x\in\mathfrak{X}} p(x) = 1$
- **By part (b) of KM-Theorem:** $\exists$ B-ary prefix code C_S with $|C_S(x)| = \ell_S(x)$ $\forall x \in \mathfrak{X}$.
  - $L_{C_S} = \sum_{x\in\mathfrak{X}} p(x)\,\ell_S(x) = \sum_{x\in\mathfrak{X}} p(x)\,\lceil -\log_B p(x)\rceil < \sum_{x\in\mathfrak{X}} p(x)\big(-\log_B p(x) + 1\big) = H_B[p] + 1$
  - in short: $L_{C_S} < H_B[p] + 1$

Symmary: Theoretical Bounds for symbol codes

EBERHARD KARLS UNIVERSITÄT TÜBINGEN

- **The Bad News:** no (uniquely decodable B-ary) symbol code can have an expected code word length smaller than the entropy $H_B[p]$ of a symbol.
- **The Good News:** one can always approach this lower bound with less than 1 bit of overhead *per symbol* (e.g., by using the *Shannon code* C_S).
- Thus, the *optimal* code C_{opt} (that minimizes L_C) satisfies:

$$\boxed{H_B[p] \leq L_{C_{\text{opt}}} < H_B[p] + 1}$$

- **Note:** The above bounds are *in expectation over all symbols* $x \in \mathfrak{X}$.
  - For any *specific* symbol $x \in \mathfrak{X}$, a code C can “violate the lower bound”: $|C(x)| < -\log_B p(x)$.
  - But: *Shannon code* satisfies $-\log_B p(x) \leq |C_S(x)| < -\log_B p(x) + 1$ *for each individual* $x \in \mathfrak{X}$.

The Source Coding Theorem [Shannon, 1948]

EBERHARD KARLS UNIVERSITÄT TÜBINGEN

- **So far:** theoretical bounds for *symbol codes*: $\boxed{H_B[p] \leq L_{C_{\text{opt}}} < H_B[p] + 1}$
- **Symbol codes are suboptimal.**
  - Always generate an *integer* number of bits per symbol.
  - Thus, overhead of up to 1 bit applies *per symbol*.

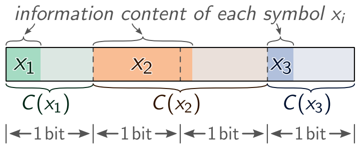

- **Practical solution:** stream codes (Lectures 5 and 6)
- **For theoretical analysis:** consider entire message $\mathbf{x} \in \mathfrak{X}^*$ as a single symbol.
  - New alphabet $\mathfrak{X}^*$ is still *countable*, thus theorems still apply.
  - Probability distribution p^* on $\mathfrak{X}^*$ can be complicated, but we'll assume it has a finite entropy $H_B[p^*] = -\sum_{\mathbf{x} \in \mathfrak{X}^*} p^*(\mathbf{x}) \log_B p^*(\mathbf{x})$.

$\Rightarrow$ The optimal uniq. dec. code C_{opt} on $\mathfrak{X}^*$ (typically *not* a symbol code on $\mathfrak{X}$) satisfies:

$$\boxed{H_B[p^*] \leq \text{expected bit rate of } C_{\text{opt}} < H_B[p^*] + 1}$$

Outlook

EBERHARD KARLS UNIVERSITÄT TÜBINGEN

- **Problem Set 2:**
  - simple examples of Shannon coding
  - entropy and information content
- **Next week:**
  - proof of optimality of Huffman coding
  - machine-learning models for lossless compression (continued in Lectures 4 and 7-9)
- **Lectures 5 & 6:** beyond symbol codes: stream codes
- **Lecture 11:** theoretical bounds for *lossy* compression ("Rate/Distortion Theory")