



Lecture 3, Part 1:

Optimality of Huffman Coding

Robert Bamler · Summer Term of 2023

These slides are part of the course “Data Compression With and Without Deep Probabilistic Models” taught at University of Tübingen. More course materials—including video recordings, lecture notes, and problem sets with solutions—are publicly available at <https://robamler.github.io/teaching/compress23/>.

Recap: Bounds for Lossless Compression



- Bounds on expected code word length of B -ary symbol codes:

$$H_B[p] \leq L_{\text{opt}} < H_B[p] + 1$$

- In addition, Shannon code C_S satisfies analogous bounds *for each symbol* $x \in \mathcal{X}$:

$$-\log_B p(x) \leq |C_S(x)| < -\log_B p(x) + 1 \quad \forall x \in \mathcal{X}$$

- Shannon code is a *near optimal* symbol code (less than 1 bit of overhead per symbol).

But how do we get an *optimal* symbol code?

Huffman Coding (recap from Lecture 1, Part 2 and Problem Set 1)



$$\begin{array}{cccc} \mathcal{X} = \{ & \text{“a”}, & \text{“b”}, & \text{“c”}, & \text{“d”} \} \\ p(x) = & 0.15 & 0.2 & 0.3 & 0.35 \end{array}$$

or

$$\begin{array}{cccc} \mathcal{X} = \{ & \text{“a”}, & \text{“b”}, & \text{“c”}, & \text{“d”} \} \\ p(x) = & 0.15 & 0.2 & 0.3 & 0.35 \end{array}$$

$$L_C = \sum_{x \in \mathcal{X}} p(x) |C(x)|$$

$$L_C = \sum_{x \in \mathcal{X}} p(x) |C(x)|$$

- ▶ **Theorem (informally):** [Huffman, 1952]
 - ▶ The Huffman algorithm constructs an optimal symbol code (i.e., it minimizes L_C).
 - ▶ If there's more than one Huffman code (due to ties) then all of them are optimal.
 - ▶ Moreover, all optimal symbol codes are equivalent to *some* Huffman code (in terms of their code word lengths $|C(x)|$).
- ▶ **Formal theorem:** assume we have:
 - ▶ finite alphabet \mathcal{X} with $|\mathcal{X}| \geq 2$
 - ▶ probability distribution $p : \mathcal{X} \rightarrow [0, 1]$ with $p(x) > 0 \forall x \in \mathcal{X}$then:
 - \forall uniquely decodable binary symbol codes $C : \mathcal{X} \rightarrow \{0, 1\}^*$ that minimize $L_C = \sum_{x \in \mathcal{X}} p(x) |C(x)|$:
 - \exists Huffman code C_H for p with $|C_H(x)| = |C(x)| \quad \forall x \in \mathcal{X}$.
- ▶ **Credits:** Our proof partially follows Jeff Miller,
https://www.youtube.com/watch?v=nvmsK__-qFg&list=PLE125425EC837021F&index=33

Lemma 1: inverse ordering

- ▶ Assume again (\star).
- ▶ Let C be an optimal prefix code for p .
- ▶ Sort the symbols by ascending probability:
$$p(x^{(1)}) \leq p(x^{(2)}) \leq p(x^{(3)}) \leq \dots \leq p(x^{(|\mathcal{X}|)})$$
break ties by code word lengths (descendingly):
 - if $p(x^{(\alpha)}) = p(x^{(\alpha+1)})$ then: $|C(x^{(\alpha)})| \geq |C(x^{(\alpha+1)})|$(break any still remaining ties arbitrarily).

then:

- (i) $|C(x^{(1)})| \geq |C(x^{(2)})| \geq |C(x^{(3)})| \geq \dots \geq |C(x^{(|\mathcal{X}|)})|$
- (ii) $|C(x^{(1)})| = |C(x^{(2)})|$

Lemma 1: inverse ordering (cont'd)

- ▶ Assume again (\star).
- ▶ Let C be an optimal prefix code for p .
- ▶ Sort the symbols by ascending probability:
$$p(x^{(1)}) \leq p(x^{(2)}) \leq p(x^{(3)}) \leq \dots \leq p(x^{(|\mathcal{X}|)})$$
break ties by code word lengths (descendingly):
 - if $p(x^{(\alpha)}) = p(x^{(\alpha+1)})$ then: $|C(x^{(\alpha)})| \geq |C(x^{(\alpha+1)})|$(break any still remaining ties arbitrarily).

then:

- (i) $|C(x^{(1)})| \geq |C(x^{(2)})| \geq |C(x^{(3)})| \geq \dots \geq |C(x^{(|\mathcal{X}|)})|$
- (ii) $|C(x^{(1)})| = |C(x^{(2)})|$

Lemma 2: weak siblings

- ▶ Assume again (\star) .
- ▶ Let C be an optimal prefix code for p .

then $\exists x, \tilde{x} \in \mathcal{X}$ with $x \neq \tilde{x}$ and:

- (i) $|C(x)| = |C(\tilde{x})| \geq |C(x')| \quad \forall x' \in \mathcal{X}$
- (ii) $C(x)$ and $C(\tilde{x})$ only differ on last bit

Proof:

- ▶ **By contradiction:** assume that such a pair does *not* exist.
- ▶ **But:** from Lemma 1, we know: the pair $(x^{(1)}, x^{(2)})$ satisfies (i)
- ▶ **Claim:** $\exists \tilde{x} \neq x^{(1)}$ such that the pair $(x^{(1)}, \tilde{x})$ satisfies both (i) and (ii).

Lemma 3: inversely ordered weak siblings

- ▶ **Recall: Lemma 1** (inverse ordering):
Among the least probable symbols, there are two symbols $x^{(1)}, x^{(2)}$ whose code words in an optimal prefix code
 - ▶ have equal length; and
 - ▶ are among the longest code words.
- ▶ **Recall: Lemma 2** (weak siblings):
Among the longest code words of an optimal symbol code, there are two code words $C(x), C(\tilde{x})$ that
 - ▶ have equal length; and
 - ▶ differ only on the last bit.
- ▶ **Note:** in general, $x^{(2)} \neq \tilde{x}$.
But: we can construct a prefix code C' with $|C'(x)| = |C(x)| \quad \forall x \in \mathcal{X}$ that satisfies both Lemma 1 and Lemma 2 for the same pair of symbols $(x^{(1)}, x^{(2)})$.



Lecture 3, Part 2:

Proof of Optimality of Huffman Coding

Robert Bamler · Summer Term of 2023

These slides are part of the course "Data Compression With and Without Deep Probabilistic Models" taught at University of Tübingen. More course materials—including video recordings, lecture notes, and problem sets with solutions—are publicly available at <https://robamler.github.io/teaching/compress23/>.

► **Assume we have:**

- finite alphabet \mathfrak{X} with $|\mathfrak{X}| \geq 2$
- probability distribution $p : \mathfrak{X} \rightarrow [0, 1]$ with $p(x) > 0 \forall x \in \mathfrak{X}$

► **Lemma 3:** assume (\star) and let C be an optimal prefix code. Then:

- ◻ prefix code C' on \mathfrak{X} with $|C'(x)| = |C(x)| \forall x \in \mathfrak{X}$, and two symbols $x^{(1)} \neq x^{(2)}$ with:
 - $C'(x^{(1)})$ and $C'(x^{(2)})$ are both longest code words, and they differ only on the last bit:
 - $p(x^{(1)})$ and $p(x^{(2)})$ have the two lowest probabilities: $p(x^{(1)}) \leq p(x^{(2)}) \leq p(x') \quad \forall x' \in \mathfrak{X} \setminus \{x^{(1)}\}$.

► **Theorem (optimality of Huffman coding):** assume (\star) . Then:

- ◻ uniquely decodable binary symbol codes $C : \mathfrak{X} \rightarrow \{0, 1\}^*$ that minimize $L_C = \sum_{x \in \mathfrak{X}} p(x) |C(x)|$:
- ◻ Huffman code C_H for p with $|C_H(x)| = |C(x)| \quad \forall x \in \mathfrak{X}$.

► **Proof:** by induction over $|\mathfrak{X}|$

- Base case ($|\mathfrak{X}| = 2$):

Induction Step (assume $|\mathfrak{X}| > 2$ and theorem holds for $|\mathfrak{X}| - 1$)

- Let C be a uniquely decodable binary symbol code on \mathfrak{X} that minimizes L_C
- Use corollary to KM-Theorem to construct a *prefix* code C' with $|C'(x)| = |C(x)| \forall x \in \mathfrak{X}$.
- Use Lemma 3 to construct a prefix code C'' with $|C''(x)| = |C'(x)| = |C(x)| \forall x \in \mathfrak{X}$ and:
- Construct the following prefix code \tilde{C} on an alphabet $\tilde{\mathfrak{X}} := (\mathfrak{X} \setminus \{x^{(1)}, x^{(2)}\}) \cup \{\square\}$:
- **Claim:** \tilde{C} is an optimal prefix code on $\tilde{\mathfrak{X}}$ (with respect to \tilde{p}).

\Rightarrow By induction hypothesis: \exists Huffman code \tilde{C}_H on $\tilde{\mathfrak{X}}$ for \tilde{p} with $|\tilde{C}_H(x)| = |\tilde{C}(x)| \quad \forall x \in \tilde{\mathfrak{X}}$.

\Rightarrow We can construct a Huffman code C_H on \mathfrak{X} for p with $|C_H(x)| = |C''(x)| = |C(x)| \quad \forall x \in \mathfrak{X}$:

So What?

You might be thinking:

"Professor, why did you just waste an hour of my life to go through a complicated proof? I would have believed you anyway."


But:

- Verification is not the point of proofs (in lectures).
- Proofs tell you:
 - *why* things are the way they are;
 - how you might be able to analyze *similar problems*.
(where you don't yet know if they're true)
- Proofs force you to think very carefully about the assumptions; this allows you to identify:
 - edge cases;
 - unnecessary assumptions (\rightarrow new applications, see Problem 3.3)

- ▶ Still widely used in practice (HTTP, zip/gzip, PNG, most JPEGs, ...)
- ▶ **But:** optimality only holds when comparing to other *symbol* codes.


Symbol codes perform poorly in the regime of low entropy per symbol.

- ▶ Consider, e.g., data source with $H_2[p] = 0.3$ bit per symbol; but $L_{CH} \geq 1$ bit per symbol.
 $\Rightarrow \sim 200\%$ overhead



- ▶ Unfortunately, this is the relevant regime for novel machine-learning based compression methods.

- ▶ **Solution:** stream codes (Lectures 5 and 6)



Lecture 3, Part 3:

Practical Compression Performance: The Modelling Gap (Kullback-Leibler Divergence)

Robert Bamler · Summer Term of 2023

These slides are part of the course "Data Compression With and Without Deep Probabilistic Models" taught at University of Tübingen. More course materials—including video recordings, lecture notes, and problem sets with solutions—are publicly available at <https://robamler.github.io/teaching/compress23/>.

Theoretical vs. Practical Bounds

- ▶ **Theoretical bounds** for an *optimal* lossless compression code: (see Lecture 2, Part 2)

$$\underbrace{H[p_{\text{data}}(\mathbf{x})]}_{\text{"entropy"}} \leq \text{expected bit rate} < H[p_{\text{data}}(\mathbf{x})] + 1$$

😊 $H[p_{\text{data}}(\mathbf{x})]$ is an intrinsic property of the *data source* (i.e., independent of any model).

☹️ We can't evaluate the *true data distribution* $p_{\text{data}}(\mathbf{x})$ for any given $\mathbf{x} \in \mathcal{X}^*$.

\Rightarrow We can't use p_{data} in an entropy coder to construct an optimal code.

\Rightarrow In fact, we can't even calculate the theoretical bound $H[p_{\text{data}}(\mathbf{x})]$.

😊 But: we can *draw samples* $\mathbf{x} \sim p_{\text{data}}$ (see next slide).

- ▶ **In practice:** (simplest case; more complicated case in Lecture 7)

1. Approximate p_{data} by some p_{model} which we *can* evaluate for all $\mathbf{x} \in \mathcal{X}^*$.
2. Optimize a compression code for p_{model} .

- ▶ **Goal:** design a compression method for some informally specified data source
 - ▶ e.g., “text that an English-speaking author might write”
 - ▶ defines the (extremely complicated) true *data generative process* with distribution p_{data} .
- ▶ **Step 1:** Collect a set \mathbb{X} of samples from the data generative process (e.g., historic books)
 - ▶ notation: $\boxed{\mathbf{x} \sim p_{\text{data}}}$ “ \mathbf{x} is sampled from the data generative process”
- ▶ **Step 2:** Create a probabilistic model p_{model} that approximates p_{data} in some way.
- ▶ **Step 3:** Use p_{model} in an entropy coder to build a (near-)optimal code C for it (and share C between sender & receiver).
 - ▶ for long messages, essentially: bit rate of code C for message $\mathbf{x} = -\log p_{\text{model}}(\mathbf{x}) \quad \forall \mathbf{x} \in \mathbb{X}^*$
- ▶ **Step 4:** In deployment, compress *new* data points $\mathbf{x} \sim p_{\text{data}}$ with C

▶ expected bit rate: $\underbrace{H(p_{\text{data}}(\mathbf{x}), p_{\text{model}}(\mathbf{x}))}_{\text{“cross entropy”}} := - \sum_{\mathbf{x} \in \mathbb{X}^*} p_{\text{data}}(\mathbf{x}) \log p_{\text{model}}(\mathbf{x}) \approx - \frac{1}{|\mathbb{X}|} \sum_{\mathbf{x} \in \mathbb{X}} \log p_{\text{model}}(\mathbf{x})$

The Modeling Gap

- ▶ **Expected bit rate in a practical setup:** $\boxed{\text{“cross entropy”} = H(p_{\text{data}}(\mathbf{x}), p_{\text{model}}(\mathbf{x}))}$
 - ▶ Motivates model training by minimizing $H(p_{\text{data}}(\mathbf{x}), p_{\text{model}}(\mathbf{x}))$ over p_{model} (\rightarrow Problem 3.2)
- ▶ **Problem 3.1:** prove that $\underbrace{H(p_{\text{data}}(\mathbf{x}), p_{\text{model}}(\mathbf{x}))}_{\text{practical bound}} \geq \underbrace{H[p_{\text{data}}(\mathbf{x})]}_{\text{theoretical bound}}$
 - ▶ equality iff $p_{\text{model}} = p_{\text{data}}$ (almost everywhere)
- ▶ **Modeling gap:** overhead (in expected bit rate) due to $p_{\text{model}} \neq p_{\text{data}}$:

$$\begin{aligned}
 D_{\text{KL}}(p_{\text{data}}(\mathbf{x}) \parallel p_{\text{model}}(\mathbf{x})) &:= H(p_{\text{data}}(\mathbf{x}), p_{\text{model}}(\mathbf{x})) - H[p_{\text{data}}(\mathbf{x})] \\
 &= \sum_{\mathbf{x} \in \mathbb{X}^*} p_{\text{data}}(\mathbf{x}) \log \frac{p_{\text{data}}(\mathbf{x})}{p_{\text{model}}(\mathbf{x})}
 \end{aligned}$$

“Kullback-Leibler divergence” aka “relative entropy”

How Good Are the Models We’ve Used So Far?

So far: $\mathbf{x} = (x_1, x_2, \dots, x_{k(\mathbf{x})})$ with some probability distribution $p_{\text{model}}(x_i)$ for all symbols x_i .

We say: symbols are modeled “i.i.d.”: *independent* and *identically distributed*.

- ▶ **identically distributed:** same distribution $p_{\text{model}}(x_i)$ for all symbols
 - ▶ Not actually necessary if we use a *prefix code*. (\rightarrow Problem 0.2 (e))
- ▶ **independent:** each symbol is modeled without regard to the other symbols.
 - ▶ Highly simplistic assumption; ignores statistical dependencies (aka *correlations*) between symbols.
 - ▶ E.g., in English text, $p_{\text{data}}(\text{‘u’})$ is much higher if the previous symbol was a ‘q’. (\rightarrow Problem 3.2)
 - ▶ Quantifying & modeling correlations requires more formal probability theory. \rightarrow **next week**

► Problem Set 3:

- prove that $D_{\text{KL}}(p \parallel q) \geq 0$
- train a machine-learning model by minimizing $H(p_{\text{data}}(\mathbf{x}), p_{\text{model}}(\mathbf{x}))$ and use it to build a compression method for written natural language

► Next week (in our regular classroom):


- probability theory
- information theoretical quantitative measure of statistical dependencies

► Afterwards: expressive probabilistic (machine-learning) models


Markov Process



Hidden Markov Model



Autoregressive Model



Latent Variable Model

